

VEXEL Feedback

What's going well (with concrete evidence)

1. Clear modular intent and separation of concerns
 - The SDK is sensibly split into focused components:
 - `src/wallet/WalletManager.ts` (key management, encryption, storage)
 - `src/signature/SignatureInjector.ts` (signing/verification)
 - `src/badge/BadgeMinter.ts` (badge minting with simulation fallback)
 - `src/haap/*` (HAAP protocol, KYC abstraction)
 - The API gateway is similarly decomposed:
 - `src/api/APIGateway.ts` wires Express + Socket.IO + middleware + Swagger

Why this matters: the “units” are small enough to test, replace, and reason about. That’s a big deal in a system that spans web2 + web3 + identity.

2. TypeScript is configured strictly
 - `tsconfig.json` has `"strict": true`. That’s the right pain now to avoid worse pain later.
3. You wrote tests for key primitives

- There are meaningful tests for the core SDK pieces:

- `src/wallet/WalletManager.test.ts`
- `src/signature/SignatureInjector.test.ts`
- `src/badge/BadgeMinter.test.ts`
- `src/haap/HAAPProtocol.test.ts`
- API middleware tests under `src/api/__tests__/`

Why this matters: you're testing the cryptographic "spine" (wallet/signing) and HAAP flow logic, which is exactly where regressions become expensive.

4. The documentation effort is unusually thorough

- You have serious phase docs and reviews (examples: `WORKFLOW.md`, `docs/SECURITY_REVIEW.md`, `docs/contracts/SECURITY_REVIEW.md`, `database/README.md`).
- The database schema is well-documented and includes triggers and views (see `database/schema.sql`, `database/README.md`).

Why this matters: docs like that are the difference between "project" and "archaeological dig".

5. Smart contract testing exists and is coherent

- Contract tests in `test/contracts/AgentHeartbeat.test.ts` cover registration, heartbeat recording, inactivity detection, and threshold updates.
-

What needs work (actionable and justified)

1. Build/package boundary is inconsistent: “features exist” but are excluded from compilation
In `tsconfig.json`, you exclude major capabilities:

- `src/database/**/*`
- `src/ipfs/**/*`
- `src/knowledge-base/**/*`
- plus `src/example.ts` and `src/service.ts`

That means those modules don’t ship in `dist/` and can’t be imported/used by consumers even though docs and examples describe them as part of the system.

Rationale: this creates a mismatch between:

- what the repo claims to provide (data layer, IPFS, knowledge migration)
 - and what the compiled package actually exports/runs.
2. Dependency hygiene is currently leaky and misleading
A quick scan of imports in build-included `src/**` shows the runtime code only needs a small set of packages (`express`, `socket.io`, `ethers`, `helmet`, etc.). But

`package.json` includes a bunch of deps that are either:

- unused by the compiled output (e.g. `redis`, `uuid`, `did-jwt`, `did-resolver`, `arweave`), or
- type packages incorrectly placed in `dependencies` (multiple `@types/*`).

Also: `dotenv` is imported by runtime code (`src/api/server.ts`) but is listed in `devDependencies`. In a real production install (where `devDependencies` aren't installed), the API server can crash.

Rationale: this inflates install surface, increases supply-chain risk, and creates “works on my machine” dependency behavior.

3. `npm test` is set up to fail in a clean environment
`jest.config.js` runs tests under `roots: ['<rootDir>/src']`, which includes:
 - integration tests that import DB/IPFS modules (`src/__tests__/repository.test.ts`, `src/__tests__/ipfs.test.ts`)

But `pg` and `kubo-rpc-client` are not declared in root `package.json`, so those imports will fail unless someone happened to install them manually.

Rationale: test runners should either:

- have all required deps declared, or
 - separate integration tests behind an explicit command so unit tests pass by default.
4. The API gateway's “agent” endpoints are stubs, so the dashboard can't actually work

`src/api/routes/agents.ts` literally returns messages like “requires database integration”.

Meanwhile, the dashboard’s `dashboard/src/services/api.ts` expects real agent objects and status fields.

Rationale: this is not a “minor missing feature”. It means one of the biggest visible deliverables (dashboard) can’t be end-to-end exercised.

5. WebSocket authentication is basically “trust me bro”

In `src/api/websocket/WebSocketServer.ts`, you accept `socket.handshake.auth.userId` as identity. No JWT validation. A client can claim to be anyone.

Rationale: if your system includes authorization logic (and it does), identity cannot be a client-supplied string.

6. Authorization is present, but not resource-scoped

`ActionVerificationMiddleware` checks roles for action types, but it does not check that:

- the human owns the target `agentId`, or
- the agent is allowed to mutate only itself

Rationale: role-based checks without ownership checks are how you end up with “any authenticated human can update any agent”.

7. Contract/subgraph mismatch

You have:

- `contracts/AgentHeartbeat.sol` (Heartbeat-based events)
- `subgraph/subgraph.yaml` configured for an `AgentRegistry` contract with events that do not exist in this repo.

Rationale: either the contract is missing, or the subgraph is for a different design. Right now it's not deployable as-is.

8. Documentation drift is already showing

Examples and docs reference flows/scripts/modules that don't line up with the shipped build and dependencies (e.g., database/IPFS/knowledge migration references in docs vs `tsconfig` excludes and missing deps).

Rationale: doc drift kills adoption faster than bugs, because it makes users distrust everything.

AI-driven task backlog (granular, with acceptance criteria)

I'm going to assume you want tasks that can be executed mechanically. Each task below has: goal, files, steps, acceptance criteria.

A) Build/package correctness (Po)

BUILD-01: Fix runtime vs dev dependency classification

Goal: runtime code must not depend on devDependencies.

Files:

- `package.json`

- `src/api/server.ts` (uses dotenv)

Steps:

1. Move dotenv from devDependencies → dependencies.
2. Run a “production-style” install test (documented): `npm ci --omit=dev` then run `node dist/api/server.js`.

Acceptance:

- API server starts without module-not-found errors when devDependencies are omitted.

BUILD-o2: Decide and enforce what the published package includes

Goal: either ship DB/IPFS/knowledge-base modules or clearly quarantine them as “internal/experimental”.

Files:

- `tsconfig.json`
- `src/index.ts`
- `docs/examples` (multiple)

Steps (Option A: ship them):

1. Remove `src/database`, `src/ipfs`, `src/knowledge-base` from `tsconfig.json` excludes.
2. Add missing runtime deps (`pg`, `kubo-rpc-client`) to `dependencies`.
3. Export the modules from `src/index.ts` (or a documented subpath export).

Acceptance:

- `npm run build` produces `dist/database/*`, `dist/ipfs/*`, `dist/knowledge-base/*`.
- Consumers can import the modules without reaching into `src/`.

Steps (Option B: don't ship them yet):

1. Keep excludes, but move those directories under something like `experimental/` or `packages/`.
2. Remove their references from “mainline” docs and root examples.

Acceptance:

- README/SETUP/TESTING do not claim those features exist in the shipped SDK.

BUILD-03: Dependency audit and pruning

Goal: remove unused deps or wire them into real implementations.

Files:

- `package.json`
- codebase import scan results

Steps:

1. Remove unused deps from `dependencies` (`redis`, `uuid`, `did-jwt`, `did-resolver`, `arweave`) unless you actually compile and ship the modules that use them.
2. Move all `@types/*` packages into `devDependencies`.

Acceptance:

- All dependencies in `dependencies` are referenced by compiled `dist/**` output (or explicitly justified in docs).

BUILD-o4: Add explicit package boundaries for multi-project repo

Goal: make dashboard/subgraph clearly separate install units.

Files:

- root `README.md`
- root `package.json`

Steps:

1. Document clearly: root package is SDK+API+contracts; dashboard and subgraph have their own `package.json`.
2. Add root scripts: `npm run dashboard:dev`, `npm run subgraph:codegen`, etc (delegating into subfolders).

Acceptance:

- New contributor can discover how to run each component from root scripts.

B) Testing & CI stability (Po)

TEST-o1: Split unit tests vs integration tests

Goal: `npm test` runs fast and passes without Postgres/IPFS.

Files:

- `jest.config.js`
- `src/__tests__/repository.test.ts`

- `src/__tests__/ipfs.test.ts`

Steps:

1. Create `jest.unit.config.js` and `jest.integration.config.js` (or use Jest “projects”).
2. Move DB/IPFS tests into `src/__tests__/integration/` and configure integration runner to include them only.
3. Update scripts:

- `test`: runs unit only
- `test:integration`: runs integration

Acceptance:

- `npm test` passes on a clean machine with no DB/IPFS.
- `npm run test:integration` runs DB/IPFS tests when services are available.

TEST-02: Fix Jest toolchain version mismatch risk

Goal: align Jest + ts-jest major versions.

Files:

- `package.json`
- `jest.config.js`

Steps:

1. Pick a consistent pair:

- Either downgrade `jest` to a version compatible with `ts-jest@29`, or
 - upgrade `ts-jest` to match Jest major version.
2. Update lockfile.

Acceptance:

- `npm test` runs without ts-jest preset errors.

CI-01: Add GitHub Actions workflow that actually runs what docs imply

Goal: enforce “green main” behavior.

Files:

- `.github/workflows/ci.yml` (new)

Steps:

1. Add jobs for:

- build (`npm ci`, `npm run build`)
- unit tests (`npm test`)
- contracts (`npm run compile`, `npm run test:contracts`)

2. Optional job: dashboard lint/test by running npm in `dashboard/`.

Acceptance:

- PRs fail if build/tests fail.
- Workflow time stays reasonable (unit vs integration split makes this realistic).

CI-02: Add integration test workflow with Postgres service

Goal: make DB schema and repository tests real, not aspirational.

Files:

- `.github/workflows/integration.yml` (new)

Steps:

1. Spin up `postgres:14` service.
2. Set env vars used by DB client.
3. Run `npm run test:integration`.

Acceptance:

- Integration workflow passes end-to-end including DB tests.

C) API gateway security and correctness (P1)

SEC-01: WebSocket authentication via JWT, not client-supplied userId

Goal: prevent identity spoofing.

Files:

- `src/api/websocket/WebSocketServer.ts`
- `src/api/middleware/auth.ts`
- `examples/websocket-client-example.ts`

Steps:

1. Require JWT in handshake (e.g., `socket.handshake.auth.token`).

2. Verify token server-side, derive userId/role from token.
3. Reject connection if missing/invalid.
4. Update example client to send token.

Acceptance:

- A client cannot connect with an invalid token.
- Connected socket sessions have server-trusted identity fields.

SEC-02: Fix CORS config contradiction

Goal: avoid invalid `origin="*"` with `credentials: true`.

Files:

- `src/api/APIGateway.ts`
- WebSocket CORS config in constructor

Steps:

1. If origin is `"*"`, set `credentials: false`.
2. If credentials are required, enforce explicit allowlist origins.

Acceptance:

- Browser clients can authenticate without CORS blocking.
- Configuration is internally consistent.

SEC-03: Resource-scoped authorization for actions

Goal: prevent cross-agent tampering.

Files:

- `src/api/middleware/actionVerification.ts`
- future DB integration in routes

Steps:

1. Extend verification to check ownership:

- If action targets `agentId`, verify request user owns `agentId` (DB lookup) or matches `agentId` for agent role.

2. Add a policy matrix: action type → required role + ownership rule.

Acceptance:

- Human users can only mutate agents they own.
- Agent role can only mutate itself.

SEC-04: Replace “token vending machine” login with a real dev/prod split

Goal: avoid shipping “POST /login, pick your role” into production by accident.

Files:

- `src/api/routes/auth.ts`

Steps:

1. Gate the current login route behind `NODE_ENV !== 'production'` or a `DEV_AUTH_ENABLED=true` flag.

2. In production mode, require real auth (even a placeholder API key check is better than none).

Acceptance:

- In production, `/api/auth/login` cannot mint admin tokens from thin air.

D) API <-> Dashboard integration (P1)

API-01: Implement real agent CRUD backed by DB repository

Goal: make `/api/agents` return actual agents, not apology messages.

Files:

- `src/api/routes/agents.ts`
- `src/database/*` (if you choose to ship/enable it)
- `database/schema.sql`

Steps:

1. Initialize `DatabaseClient` in API gateway (configurable via env).
2. Use `AgentRepository` to implement:
 - `GET /api/agents`
 - `GET /api/agents/:agentId`
 - `PUT /api/agents/:agentId/status`
 - `GET /api/agents/:agentId/capabilities`

3. Normalize response shapes (don't nest `agents` inside `data.agents` while dashboard expects `data` array, etc).

Acceptance:

- Dashboard `getAgents()` returns a list of real agents.
- Status updates persist and can be re-fetched.

API-02: Add request validation (zod or similar)

Goal: stop trusting `req.body` to behave like an adult.

Files:

- `src/api/routes/auth.ts`
- `src/api/routes/agents.ts`

Steps:

1. Add schemas for each route input.
2. Validate and return 400 with specific errors.

Acceptance:

- Malformed inputs get predictable 400 responses.
- No route relies on implicit types from JSON bodies.

API-03: Make Swagger generation work from dist

Goal: `/api-docs` shouldn't silently break after compilation.

Files:

- `src/api/APIGateway.ts`

Steps:

1. Change `apis: ['./src/api/routes/*.ts']` to a path that exists at runtime:

- use `path.join(__dirname, 'routes/*.js')` in `dist`
- or include source files intentionally

Acceptance:

- Swagger UI shows endpoints when running compiled `dist` server.

DASH-01: Align dashboard agent types with API reality

Goal: stop type drift (active/inactive vs ACTIVE/SLEEP/TERMINATED).

Files:

- `dashboard/src/types/index.ts`
- `dashboard/src/services/api.ts`

Steps:

1. Either:

- map API runtime status enum → dashboard status union, or
- change dashboard to use API enum directly.

2. Update `getDashboardStats()` logic accordingly.

Acceptance:

- Dashboard shows correct status counts from real API responses.

E) Wallet/key security hardening (P1)

WALLET-01: Enforce encryption key in production and fix file permissions

Goal: wallet storage isn't "security theater".

Files:

- `src/wallet/WalletManager.ts`

Steps:

1. If `NODE_ENV === 'production'` and no `WALLET_ENCRYPTION_KEY`, throw at startup (not later).
2. When writing wallet files, set permissions `0o600`.

Acceptance:

- Production run fails fast without encryption key.
- Wallet files are not world-readable on disk.

WALLET-02: Stop returning mnemonic by default

Goal: reduce accidental secret exfiltration.

Files:

- `src/wallet/WalletManager.ts`
- `src/wallet/WalletManager.test.ts`

Steps:

1. Add option `returnMnemonic?: boolean` default false.

2. Only include mnemonic when explicitly requested.

Acceptance:

- `createWallet()` does not leak mnemonic unless configured.

F) HAAP protocol correctness & durability (P2)

HAAP-o1: Replace weak token ID generation with crypto-grade randomness

Goal: avoid predictable token IDs.

Files:

- `src/haap/HAAPProtocol.ts`

Steps:

1. Generate token IDs via `crypto.randomBytes(16).toString('hex')` or UUID v4.

Acceptance:

- Token IDs are not based on time/random substring.

HAAP-o2: Persist tokens and KYC status (or remove the pretense)

Goal: tokens surviving process restarts is table-stakes.

Files:

- `src/haap/HAAPProtocol.ts`

- `src/haap/KYCSERVICE.ts`

Steps (choose one):

1. Implement Redis-backed storage (you already depend on redis but don't use it).
2. Or implement DB-backed storage.

Acceptance:

- After restart, `validateToken(tokenId)` still works.

G) Contracts + subgraph coherence (P2)

CHAIN-01: Fix DID squatting / unauthorized registration in AgentHeartbeat

Goal: prevent attackers registering someone else's DID.

Files:

- `contracts/AgentHeartbeat.sol`
- `test/contracts/AgentHeartbeat.test.ts`

Steps:

1. Require `msg.sender == agentAddress` in `registerAgent`, or require an EIP-712 signature proving control.
2. Add tests for attempted squatting.

Acceptance:

- Unauthorized registration attempts revert.
- Tests cover the attack path.

SUBGRAPH-01: Make subgraph match the deployed contract set

Goal: subgraph should be deployable.

Files:

- `subgraph/subgraph.yaml`

- `subgraph/abis/*`

- `subgraph/src/mapping.ts`

Steps (pick one direction):

1. If `AgentRegistry` is the intended contract, add `contracts/AgentRegistry.sol` + hardhat compile artifacts + deployment scripts.
2. If `AgentHeartbeat` is intended, rewrite subgraph to index its events and update schema accordingly.

Acceptance:

- `graph codegen` and `graph build` succeed.
- Indexed events match actual chain events.

H) Documentation drift control (P2)

DOC-01: Create a “single source of truth” run matrix

Goal: stop docs lying by accident.

Files:

- `README.md`
- `SETUP.md`
- `TESTING.md`

Steps:

1. Add a matrix: component → where to run → command → required services.
2. Remove/flag any steps that reference excluded modules unless those modules are shipped.

Acceptance:

- A new dev can run SDK tests, API server, dashboard, and contract tests without guessing.

DOC-02: Make examples compile against the published surface

Goal: examples should only import what's actually exported/built.

Files:

- `examples/*`
- `src/index.ts`

Steps:

1. For each example, ensure imports target the public API.
2. Remove examples that depend on excluded modules, or move them under `experimental/`.

Acceptance:

- `npx ts-node examples/<file>.ts` works after `npm install`.
-